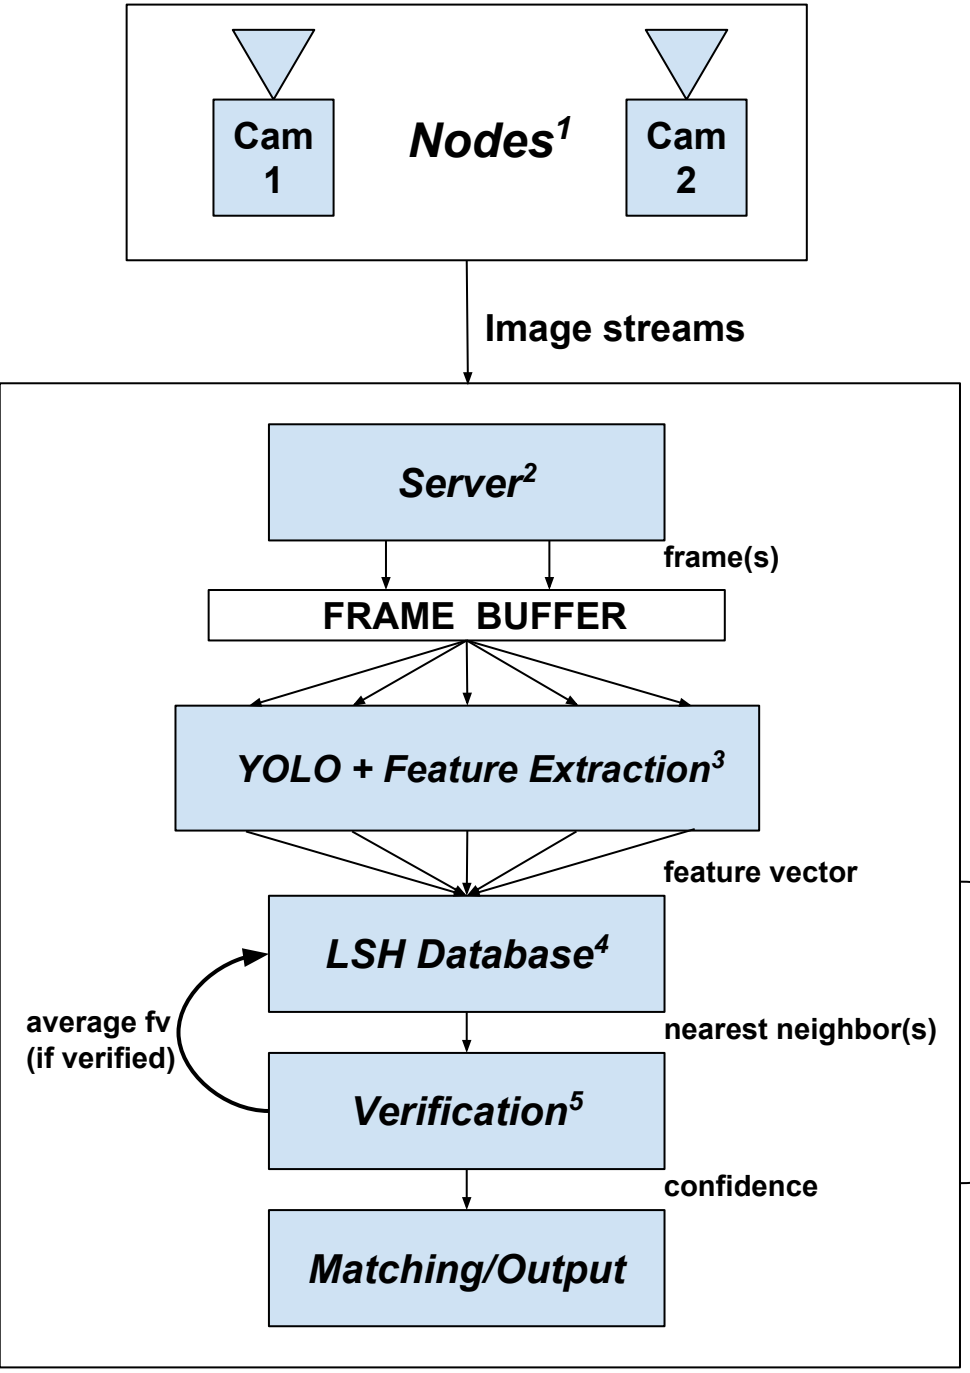

Potential Issues:

1 Pedestrian Detection Trigger (HOG+SVM - too slow. Currently using BG subtraction)

2 How many frames to "detect" from 1st camera?
How many to use from 2nd camera to re-id? (both have an fps of 7)

3 Find ways to optimize YOLO+FE
Check how the feature vectors of the extraction look compared to the Verification's feature vectors

4 How accurate is averaging the old and new feature vectors?
How accurate is LSH? Is discretized and/or PCA more accurate than RBPT?

5 Compare Verification's feature vectors to Feature Extraction's (same as 4)

If verification matches on frames from *same* camera, then compute average of original and newly extracted feature vectors and update the object. If it matches on different cameras, then the pedestrian has been re-identified